

CS 109b – Introduction to Bayesian methods

Bayesian statistics:

- Foundations and philosophy of Bayesian statistics
- Recipe to performing a Bayesian analysis
- Advantages and disadvantages to the Bayesian paradigm
- Modern approach to Bayesian statistics
- Example analyses

Are you a Bayesian?

In the following situations, what would you conclude about the (future) probability of success?

1. A music expert claims he can identify whether a piece of music was written by Mozart versus Haydn.

In 10 trials, the expert identifies the composer correctly 10 times.

2. A drunken friend claims he can predict the result of coin flips.

In 10 trials, he answers correctly 10 times.

3. A coffee connoisseur claims to be able to tell the difference between whether cream is poured in coffee first, or vice versa.

In 10 trials, the person answers correctly 10 times.

If you drew different conclusions in all three situations, then you are likely already Bayesian.

- One key feature of Bayesian statistics is that it is a “learning paradigm.” You cannot (and should not) ignore knowledge you currently have in summarizing conclusions based on data.
- In the three example situations, we may have had strong prior beliefs that influenced our conclusions about the data. In classical settings, it is not straightforward to incorporate such beliefs.
- Bayesian statistics provides a formal way to incorporate prior information into an analysis.

Question:

What is the underlying difference between classical and Bayesian statistics?

Answer:

The definition of probability.

Example: Flip a fair coin; the probability the coin lands heads is $\frac{1}{2}$.

What does “the probability is $\frac{1}{2}$ ” mean?

Classical interpretation:

On repeated flips of the coin, the coin lands heads $\frac{1}{2}$ of the time in the long run.

Frequency definition of probability: Probability of an event is its **long-run frequency** of occurrence.

This definition is useful for describing the likelihood of potentially observable data.

Bayesian interpretation:

There is an equal probability of heads and tails due to the symmetry of the coin, so that $\frac{1}{2}$ is an assessment of my belief that the coin will land heads.

Subjective definition of probability: Probability of an event is one’s **degree of belief** that the event will occur (1 means that it is certain to occur, and 0 means that it is certain not to occur).

This definition is useful for quantifying beliefs about non-data events (but also data events).

Classical statistics, which is founded on the frequency definition of probability, is usually called Frequentist statistics. Much of the theory developed by Jerzy Neyman in early 1900s.

Implications to a theory of statistical inference:

- Statistics founded on the Frequentist definition of probability can only assign probabilities to future data or potentially observable quantities.
- Statistics founded on the subjective definition of probability can consider probabilities of values of unknown parameters (as well as values of potentially observable quantities).

Variation of coin flipping:

Suppose yesterday, at 5:00pm, standing at home, I flipped a fair coin and recorded whether it landed heads.

What is the probability it landed heads?

Answer:

If you are a Bayesian, the answer is unequivocally $\frac{1}{2}$.

If you are a Frequentist, you are in a bit of a bind because I am not asking about a long-run frequency – I am asking about what happened specifically at 5:00pm yesterday.

A Frequentist would need to say either

- that the probability is either 0 or 1, but would not know which is correct, or
- need to resort to frequencies of events across different universes within a multiverse!

Because the goal of statistical inference is to infer unknown quantities from data, Frequentists often find themselves in awkward philosophical quandaries.

Example: Confidence intervals

Suppose y_1, \dots, y_n are a random sample of n values from $N(\mu, \sigma^2)$.

Formula for a 95% confidence interval is

$$(\bar{y} - t_{n-1}^* \frac{s}{\sqrt{n}}, \bar{y} + t_{n-1}^* \frac{s}{\sqrt{n}}).$$

where \bar{y} is the sample mean, s is the sample standard deviation, and t_{n-1}^* is a critical value from the t -distribution on $n - 1$ degrees of freedom.

Before observing data, 95% of the time these endpoints will contain μ , the value we are trying to infer.

- Notice that the Frequentist definition applies to characterizing the **procedure**. That is, if we apply the procedure for determining 95% confidence intervals according to the formula, we will be correct 95% of the time.
- Once data are observed, the coverage probability is either 0 or 1.

Here is a dialog I once heard between a Frequentist and his employer where the Frequentist tried to explain confidence intervals.

A Conversation between a Frequentist Consultant and his insightful Employer

Employer: So, do you have the results of the drug trial analyses?

Frequentist: Yes, I do.

Employer: Let's hear it.

Frequentist: Okay. I computed a 95% confidence interval for the average decrease in diastolic blood pressure from taking the medication, and I got (10.18, 14.92) mm Hg.

Employer: Great! What does that mean?

Frequentist: It means we should be 95% "confident" that the true mean decrease is between 10.18 and 14.92.

Employer: Oh, you mean that there is a 95% probability that the mean decrease is between 10.18 and 14.92 mm Hg. Got it!

Frequentist: Well, not exactly.

Employer: Huh?

Frequentist: It means that if we were to compute “95% confidence intervals” from many data samples, about 95% of them would contain the true mean decrease in blood pressure.

Employer: Um... But what about **this** data set?

Frequentist: I don’t know.

Employer: What do you mean you don’t know?

Frequentist: All I know is that in 95% of the analyses I perform, my 95% confidence intervals contain the parameter I’m estimating.

Employer: But I’ve hired you to draw conclusions from this data set!

Frequentist: Well, you can hope that your data set is “typical.”

Employer: What are you talking about? What does that have to do with 95%?

Frequentist: If your data set was one of the 95% “typical” data sets you would observe, then the confidence interval would contain the mean decrease in blood pressure.

Employer: But what does that tell me about the confidence interval you reported to me?

Frequentist: Nothing.

Employer: Nothing?

Frequentist: Can I get my consulting fee? I have to go.

Employer: No way! You’re FIRED!!!

Bayesians compute and report intervals as well, but they have different meaning because they are based on subjective probability.

- For example, a Bayesian can compute an interval in which there is 95% probability (as a degree of belief) that μ is in the interval.
- Note that in this case the endpoints of the interval are values computed from data.

Here’s a conversation between the same employer, but with a Bayesian.

A Conversation between a Bayesian Consultant and his insightful Employer

Employer: So, do you have the results of the drug trial analyses?

Bayesian: Yes, I do.

Employer: Let's hear it.

Bayesian: Okay. I computed a 95% central posterior interval for the average decrease in diastolic blood pressure from taking the medication, and I got (10.04, 14.83) mm Hg.

Employer: Great! What does that mean?

Bayesian: It means we should be 95% certain that the true mean decrease is between 10.04 and 14.83 mm Hg.

Employer: Oh, you mean that there is a 95% probability that the mean decrease is between 10.04 and 14.83. Got it!

Bayesian: Exactly!

THE END!

Simple but practical example: Back to the coffee connoisseur

- Suppose p is the probability the coffee connoisseur correctly guesses whether cream is poured in coffee first, or vice versa.
- We observed 10 correct guesses out of 10 tries.
- The classical point estimate of p is the sample proportion, that is $\hat{p} = 10/10 = 1$ (with a standard error of 0).

This means that our best guess is that the coffee connoisseur is 100% accurate.

How can that be?

Bayes theorem (or Bayes Rule):

Before introducing the Bayesian approach to statistics, it is worthwhile reviewing Bayes theorem for discrete events.

Suppose A and B are two events where A temporally or causally comes before B .

Example: A is the event that you attend this lecture or watch the video, and B is the event that you score a 5 on your Bayesian homework.

The use of Bayes Rule is to determine the probability of a “past” or “cause” event, given a current event has happened.

More formally, we use Bayes Rule to obtain $P(A|B)$. This is a “detection” probability – given what we now know (B), what is the probability of a suspected event in the past (A)?

This is given by Bayes Rule:

$$\Pr(A|B) = \frac{\Pr(A) \Pr(B|A)}{\Pr(A) \Pr(B|A) + \Pr(A^c) \Pr(B|A^c)}$$

Usually $\Pr(A)$, $\Pr(B|A)$ and $\Pr(B|A^c)$ can be determined easily.

Example: Two coins, one fair and one double-headed, are placed in a box. One coin is chosen at random, and flipped. If it lands heads, what is the probability it is the double-headed coin?

Answer: Let

A = The double-headed coin is selected

B = The coin selected lands heads

Can write down

$$\begin{aligned}\Pr(A) &= 1/2 \\ \Pr(B|A) &= 1 \\ \Pr(B|A^c) &= 1/2\end{aligned}$$

So, by Bayes Rule,

$$\begin{aligned}\Pr(A|B) &= \frac{\Pr(A) \Pr(B|A)}{\Pr(A) \Pr(B|A) + \Pr(A^c) \Pr(B|A^c)} \\ &= \frac{(1/2)(1)}{(1/2)(1) + (1/2)(1/2)} = \frac{2}{3}\end{aligned}$$

The probability that the double-headed coin was selected (the past event) given that the coin landed heads (the current event) is $\frac{2}{3}$.

Application of Bayes Rule: Naive Bayes Spam Filter

Goal: Want to develop a way to assess the probability a new incoming e-mail is spam.

Assumptions:

- You have a large collection of e-mails to develop your method. From these e-mails, you can easily characterize as “spam” versus “not spam.”
- You have a program that can search for occurrences of words in your e-mails.

Suppose we identify 10 word phrases common in spam. For example,

- earn money, click to remove, visit our website, stay in shape, act now, free membership

Let W_j be the event that the j -th phrase appears in the e-mail. Let W_j^c (complement of W_j) denote the event that the j -th phrase is not in the e-mail.

Assume that an e-mail contains the 1st and 10th words, but none of the others. Want to compute or estimate

$$\Pr(\text{spam} \mid W_1, W_2^c, \dots, W_9^c, W_{10}).$$

By Bayes Rule,

$$\Pr(\text{spam} \mid W_1, W_2^c, \dots, W_9^c, W_{10}) = \frac{\Pr(W_1, W_2^c, \dots, W_9^c, W_{10} \mid \text{spam}) \Pr(\text{spam})}{\Pr(W_1, W_2^c, \dots, W_9^c, W_{10} \mid \text{spam}) \Pr(\text{spam}) + \Pr(W_1, W_2^c, \dots, W_9^c, W_{10} \mid \text{spam}^c) \Pr(\text{spam}^c)}.$$

We can empirically estimate $\Pr(\text{spam})$ from the sample proportion of spam you receive.

Also, $\Pr(\text{spam}^c) = 1 - \Pr(\text{spam})$.

Estimating

$$\Pr(W_1, W_2^c, \dots, W_9^c, W_{10} \mid \text{spam})$$

and

$$\Pr(W_1, W_2^c, \dots, W_9^c, W_{10} \mid \text{spam}^c)$$

can also in principle be accomplished empirically by finding the sample proportion of e-mails in your development sample with the 1st and 10th phrase in the note, but none of the others.

Problem: May be very few e-mails with exactly the 1st and 10th phrase in the note.

This problem is even worse if we considered, say, 100 spam phrases.

Possible solution: Naive Bayes

Use the following approximation:

$$\Pr(W_1, W_2^c, \dots, W_9^c, W_{10} \mid \text{spam}) \approx \Pr(W_1 \mid \text{spam}) \Pr(W_2^c \mid \text{spam}) \cdots \Pr(W_9^c \mid \text{spam}) \Pr(W_{10} \mid \text{spam})$$

$$\Pr(W_1, W_2^c, \dots, W_9^c, W_{10} \mid \text{spam}^c) \approx \Pr(W_1 \mid \text{spam}^c) \Pr(W_2^c \mid \text{spam}^c) \cdots \Pr(W_9^c \mid \text{spam}^c) \Pr(W_{10} \mid \text{spam}^c)$$

This approximation assumes that the 10 phrase occurrences are *conditionally independent* given the spam status of each e-mail.

This is a huge simplification because it ignores the possibility that some spam phrases tend to travel together.

But this approach has the following advantages:

- Estimates of $\Pr(W_j \mid \text{spam})$ and $\Pr(W_j \mid \text{spam}^c)$ are likely to be accurate.
- Easy to compute.
- If the occurrences of pairs, triples, etc., of spam phrases are not very highly dependent, assuming conditional independence can give accurate estimates of combinations of phrases that are not observed in the development sample of e-mails.

See <http://homes.cs.washington.edu/~pedrod/papers/cacml2.pdf> to see how Naive Bayes can outperform state-of-the-art learner.

From Bayes Rule to Bayesian statistics:

When we set up probability models for data, we specify

- Unknown model parameters
- How data are generated based on the model parameters

Can think of the model parameters coming first, and data coming second (based on the values of the model parameters).

Bayes theorem provides the tool to compute the probability of model parameter values based on the observed data.

In Bayesian statistics, all quantities (data, unknown parameters) have probability distributions to describe uncertainty.

Parameters: Uncertainty in model parameters are described through probability distributions $p(\theta)$. For instance, we may believe *a priori* that the mean adult male weight is uniformly distributed from 160 lbs to 200 lbs.

Data: Observations obtained from a data-generating mechanism are assumed to follow a probability distribution $p(y|\theta)$. For example, we may assume adult male weights are normally distributed around an unknown mean.

Note that assuming a probability distribution on a parameter does not mean that the parameter is “random.” It is actually a fixed value, but our uncertainty about it is represented in the form of a probability distribution.

Summary of the Bayesian approach:

A typical Bayesian analysis can be outlined in the following steps.

1. Formulate a probability model for the data.
2. Decide on a prior distribution for the unknown model parameters.
3. Observe the data, and construct the likelihood function based on the data.
4. Determine the posterior distribution.
5. Summarize important features of the posterior distribution, or calculate quantities of interest based on the posterior distribution.

Keeping your eye on the ball:

The main goal of Bayesian inference is to obtain the **posterior distribution** of the unknown parameters. This is the probability distribution that describes the state of knowledge about the parameters once the data have been observed.

After the posterior distribution has been determined, just about any inferential question can be answered.

Simple example to illustrate steps in Bayesian analysis:

Suppose we are examining patterns of 30-day mortality of heart attack patients at a local hospital.

- Among 5 randomly selected heart attack patients, 1 died and 4 survived beyond 30 days.
- Let θ be the probability a patient dies within 30 days after admission for a heart attack.
- We want to make inferences about θ from the data.

First step: Formulate the probability model

For $i = 1, \dots, 5$, let

$$Y_i = \begin{cases} 0 & \text{if the } i\text{-th patient survives 30 days} \\ 1 & \text{if the } i\text{-th patient dies within 30 days} \end{cases}$$

Probability model for the Y_i :

$$\Pr(Y_i = y \mid \theta) = \begin{cases} \theta & \text{for } y = 1 \\ 1 - \theta & \text{for } y = 0 \end{cases}$$

More compactly,

$$P(Y_i = y \mid \theta) = p(y \mid \theta) = \theta^y (1 - \theta)^{1-y},$$

for $y = 0, 1$. This is a Bernoulli model for the data.

Second step: Choosing a prior distribution


The prior distribution of the unknown parameter(s) represents the state of knowledge prior to observing the data.

For this setting, let us assume that all values of θ between 0 and 1 are equally believable.

Formally, this translates to

$$p(\theta) = 1$$

for $0 \leq \theta \leq 1$.



Third step: Construct the likelihood function


The likelihood function is the probability of the data (probability density, for continuous outcomes) conditional on the parameter(s), viewed as a function of the parameter(s).

$$p(y_1, y_2, y_3, y_4, y_5 \mid \theta) = \prod_{i=1}^5 \theta^{y_i} (1 - \theta)^{1-y_i} = \theta(1 - \theta)^4.$$

Thus, the likelihood function is

$$L(\theta \mid \mathbf{y}) = \theta(1 - \theta)^4,$$

where $\mathbf{y} = (y_1, y_2, \dots, y_5)$ are the five binary outcomes.



Fourth step: Determine the posterior distribution

The posterior distribution is the probability distribution of the unknown parameter(s) conditional on the observed data.

This is determined by Bayes rule:

$$p(\theta | \mathbf{y}) = \frac{p(\theta)p(\mathbf{y} | \theta)}{\int p(\theta)p(\mathbf{y} | \theta) d\theta} = \frac{p(\theta)p(\mathbf{y} | \theta)}{p(\mathbf{y})}$$

$$\propto p(\theta)p(\mathbf{y} | \theta) \propto p(\theta)L(\theta | \mathbf{y})$$

where “ \propto ” means “is proportional to.”

The key tool in Bayesian statistics:

$$p(\theta | y) \propto p(\theta)L(\theta | y)$$

Posterior \propto Prior \times Likelihood

Applying the key tool in our setting:

We have

$$\begin{aligned}p(\theta) &= 1 \\L(\theta \mid \mathbf{y}) &= \theta(1 - \theta)^4\end{aligned}$$

This means

$$p(\theta \mid \mathbf{y}) \propto p(\theta)L(\theta \mid \mathbf{y}) = (1)\theta(1 - \theta)^4 = \theta(1 - \theta)^4$$

This in turn means that

$$p(\theta \mid \mathbf{y}) = c\theta(1 - \theta)^4$$

for some constant c (that does not depend on θ).

It turns out

- $p(\theta) = 30 \theta(1 - \theta)^4$ integrates to 1, so that $c = 30$ in this case.
- This is the density for a Beta distribution with parameters 2 and 5 (one more than the exponents of θ and $(1 - \theta)$, respectively), abbreviated


$$\theta \sim \text{Be}(2, 5).$$

This is a probability distribution used frequently by Bayesians.

For general Beta distributions, write $\text{Be}(\alpha, \beta)$.

Some properties of the Beta distribution:

$$\begin{aligned}\text{E}(\theta \mid \alpha, \beta) &= \frac{\alpha}{\alpha + \beta} \\ \text{Var}(\theta \mid \alpha, \beta) &= \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)} \\ \text{Mode}(\theta \mid \alpha, \beta) &= \frac{\alpha - 1}{\alpha + \beta - 2}\end{aligned}$$



Fifth step: Summarizing the posterior distribution

- Posterior mean, $E(\theta|\mathbf{y})$.
- Posterior mode (value of θ that maximizes $p(\theta | \mathbf{y})$).
- Central posterior interval for θ .
- Highest posterior density (HPD) region for θ – shortest interval with specified probability. Usually more difficult to compute than central intervals.


Posterior summaries for our setting:

Based on properties of the Beta distribution

- $E(\theta|\mathbf{y}) = \frac{2}{2+5} \approx 0.286$
- $\text{mode}(\theta|\mathbf{y}) = \frac{2-1}{2+5-2} = 0.2$
- 95% central posterior interval: Can compute the 2.5%-ile and 97.5%-ile numerically in R.

(0.0433, 0.641)

Compare to the Frequentist 95% confidence interval (0.011, 0.701).



Comments:

- The Bayesian approach recognizes the asymmetry in inferences about θ , whereas the standard Frequentist approach does not.
- The same approach outlined above applies identically to multi-parameter models.
- Frequentist approach can be improved, but need to use fancier machinery, e.g., the bootstrap, but even that is not often reliable with small samples.

Bayesian inference as a learning model:

The Bayesian approach facilitates the idea that one learns about the state of the world as data are accumulated.

Sequential nature of Bayesian analysis:

“Yesterday’s posterior is today’s prior”

Instead of performing a single analysis with the 5 observations, I could have performed (up to) 5

analyses in sequence:

Suppose $y_1 = 0, y_2 = 0, y_3 = 1, y_4 = 0$, and $y_5 = 0$.


Start with $p(\theta)$.

Observe $y_1 = 0$, and determine $p(\theta \mid y_1 = 0)$.

Now use $p(\theta \mid y_1)$ as the prior distribution for the second observation.

Observe $y_2 = 0$, and determine $p(\theta \mid y_1 = 0, y_2 = 0)$.

And so on! You will obtain the same posterior distribution as analyzing the data simultaneously.



Application: Inferring authorship

In 1963, a classic paper by Mosteller and Wallace was published that was one of the first large-scale applications of Bayesian statistics.

- The *Federalist* papers were published anonymously in 1787-1788 by Alexander Hamilton, John Jay, and James Madison to convince the State of NY to ratify the constitution.
- 62 of the essays are of known authorship, and 3 were jointly written.
- The authorship of the remaining 12 is in dispute between Hamilton and Madison.

Idea: Use the frequency of “filler” words to discriminate between Hamilton and Madison.

Use of “upon” by Hamilton versus Madison in 48 Hamilton and 50 Madison papers (not just Federalist papers):

Rate of use per 1000 words	Authored by Hamilton	Authored by Madison	12 Disputed Papers
Exactly 0	0	41	11
(0.0, 0.4)	0	2	0
[0.4, 0.8)	0	4	0
[0.8, 1.2)	2	1	1
[1.2, 1.6)	3	2	0
[1.6, 2.0)	6	0	0
[2.0, 3.0)	11	0	0
[3.0, 4.0)	11	0	0
[4.0, 5.0)	10	0	0
[5.0, 6.0)	3	0	0
[6.0, 7.0)	1	0	0
[7.0, 8.0)	1	0	0
Total	48	50	12

To the eye, clear evidence that the 12 disputed papers are most like Madison. How can Bayesian statistics help?

Details in the context of musical authorship:

Who wrote this song?

If you said *The Beatles*, you are right.

But which Beatle?

- As members of the Beatles, John Lennon and Paul McCartney agreed to credit songs written by each as “Lennon/McCartney.”
- In most cases, the actual composer of the song is known, but for a small sample of their songs the authorship is in dispute.
- *In My Life* is one song that is in dispute.

Bayesian analysis of Beatles authorship:

Data:

- All Lennon/McCartney songs during 1962-1966 (from *Please Please Me* to *Revolver*). Total of 70 songs with known authorship, and several with unknown authorship.
- For each song, decompose the melodic line into a sequence of *melodic note transition couplets*.

Suppose that there are K possible transition couplets.

Let

$$(y_1^L, y_2^L, \dots, y_K^L)$$

be the counts of the couplets across 39 Lennon songs, and

$$(y_1^M, y_2^M, \dots, y_K^M)$$

be the counts of the couplets across 31 McCartney songs.

Let $N^L = \sum_{k=1}^K y_k^L$ be the total number of Lennon couplets and $N^M = \sum_{k=1}^K y_k^M$ be the total number of McCartney couplets.

The model we assume is

$$\begin{aligned} p(y_1^L, y_2^L, \dots, y_K^L | \theta^L) &\propto (\theta_1^L)^{y_1^L} (\theta_2^L)^{y_2^L} \dots (\theta_K^L)^{y_K^L} \\ p(y_1^M, y_2^M, \dots, y_K^M | \theta^M) &\propto (\theta_1^M)^{y_1^M} (\theta_2^M)^{y_2^M} \dots (\theta_K^M)^{y_K^M} \end{aligned}$$

which is a Multinomial model.

We also assume a Dirichlet prior distribution for the θ_k^L and θ_k^M :

$$\begin{aligned} p(\theta_1^L, \theta_2^L, \dots, \theta_K^L) &\propto (\theta_1^L)^\alpha (\theta_2^L)^\alpha \dots (\theta_K^L)^\alpha \\ p(\theta_1^M, \theta_2^M, \dots, \theta_K^M) &\propto (\theta_1^M)^\alpha (\theta_2^M)^\alpha \dots (\theta_K^M)^\alpha \end{aligned}$$

with α set in advance (though can be estimated via cross-validation).

Suppose a song of unknown authorship has melodic transition couplet frequencies (y_1^0, \dots, y_K^0) , with $N^0 = \sum_{k=1}^K y_k^0$.

Let R_L and R_M be the so-called “marginal likelihoods” under the assumption that the song of unknown authorship is Lennon versus McCartney.

Can show

$$\begin{aligned} R_L = & \left(\frac{N^0! N^L!}{(N^0 + N^L)!} \right) \left(\frac{\prod_{k=1}^K (y_k^0 + y_k^L)!}{\prod_{k=1}^K y_k^0! \prod_{k=1}^K y_k^L!} \right) \times \\ & \left(\frac{(N^0 + N^L)! \Gamma(K\alpha)}{\Gamma(N^0 + N^L + K\alpha)} \prod_{k=1}^K \frac{\Gamma(y_k^0 + y_k^L + \alpha)}{(y_k^0 + y_k^L)! \Gamma(\alpha)} \right) \left(\frac{(N^M)! \Gamma(K\alpha)}{\Gamma(N^M + K\alpha)} \prod_{k=1}^K \frac{\Gamma(y_k^M + \alpha)}{(y_k^M)! \Gamma(\alpha)} \right). \end{aligned}$$

The analogous expression can be computed for R_M switching the roles in the above equation for L and M .

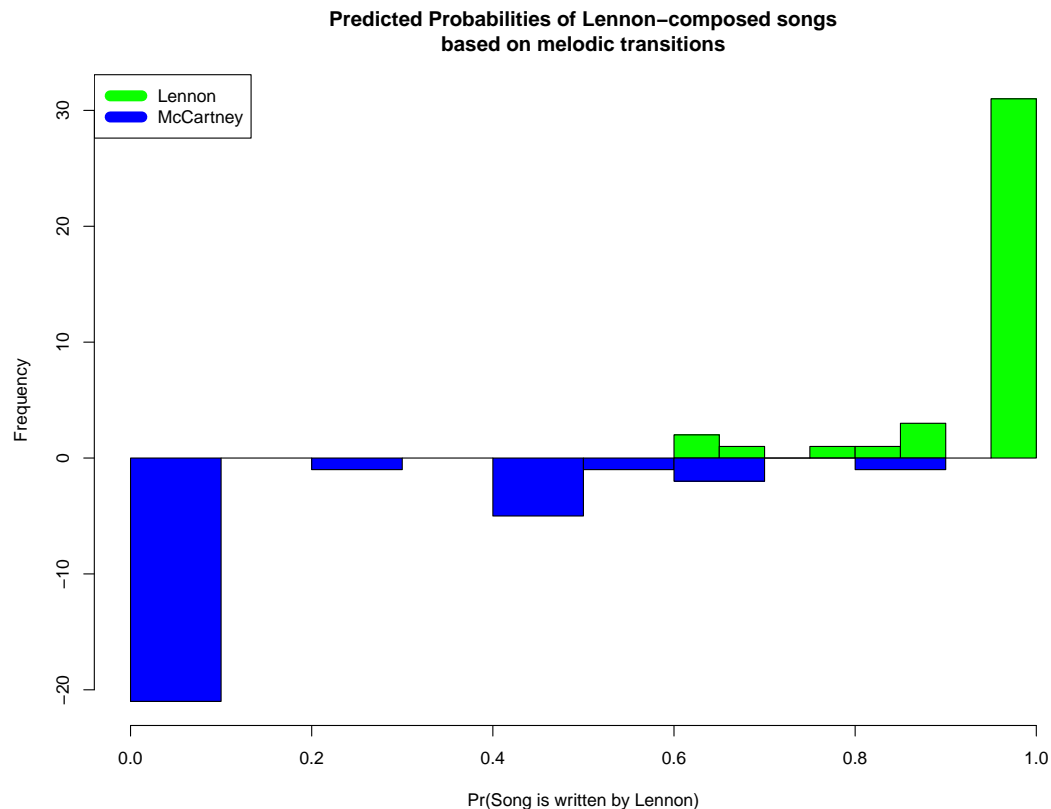
The posterior probability that the new song is written by Lennon is

$$\Pr(\text{Lennon} \mid \text{data}) = \frac{R_L}{R_L + R_M}.$$

With α optimized as 5.53×10^{-5} through cross-validation, we get

Song	Pr(Lennon song)
Misery	0.999
Ask Me Why	0.999
Do You Want to Know a Secret	0.726
Wait	0.670
In My Life	0.585

Inconclusive for “In My Life”!



Modern Bayesian approach to data analyses:

The previous discussion laid out the foundations of the Bayesian approach, but more complex modeling situations can present great challenges.

- The posterior density for most real-life models can be time-consuming to determine.
- Even if one can write down the posterior density, it is almost always the case that it is difficult to summarize using standard analytic tools.

The modern approach: Summarizing posterior distributions through Monte Carlo simulation

Monte Carlo simulation:

The posterior distribution is a probability distribution, and the goal of a Bayesian analysis is to study the posterior distribution, or derive summaries from the posterior distribution.

Summarizing computer-simulated values from the posterior distribution is a legitimate alternative to analytic summaries.

Monte Carlo summaries of probability distributions:

- Rather than go through laborious calculations to obtain the mean, median, variance, percentiles, etc., of a probability distribution, one can calculate the sample-version of the distribution from the simulated sample.
- Usually need to simulate a very large sample in order to precisely approximate the generating distribution.

Example: Summarizing $\text{Be}(2, 5)$ distribution

Rather than report the analytically-determined summaries of the Beta posterior distribution for the 30-day mortality example, do the following:

1. Simulate 10,000 values from $\text{Be}(2, 5)$.
2. Report *sample* summaries from the distribution of 10,000 simulated values.

```
> theta = rbeta(10000, 2, 5)
> head(theta) # first six simulated values
[1] 0.2244411 0.4130596 0.2961517 0.2133488 0.1883770 0.6113721
> mean(theta) # sample mean
[1] 0.2868297
> quantile(theta, probs=c(0.025, 0.975)) # sample quantiles
      2.5%      97.5%
0.04411736 0.64922037
```

Recall that $E(\theta | \alpha = 2, \beta = 5) = 0.286$ and the 95% central posterior interval was (0.0433, 0.641).

Predictive inference:

One of the real strengths of the Bayesian setting is the ability to produce model-based predictions that accounts for the uncertainty in parameter inferences.

Goal: Want to determine $p(\tilde{y} | \mathbf{y})$, that is, the probability distribution for a new value \tilde{y} given the data we already analyzed. This distribution is called the **posterior predictive distribution**.

Notice that θ is no longer in the expression – the posterior predictive distribution “averages out” the uncertainty about θ .

Can show analytically using tools of calculus that

$$p(\tilde{y} \mid \mathbf{y}) = \begin{cases} 5/7 & \text{if } \tilde{y} = 0 \\ 2/7 & \text{if } \tilde{y} = 1 \end{cases}$$

Thus we can conclude that there is a 2/7 chance that a new (randomly selected) patient entering the hospital with a heart attack would die within 30 days.

Computing $p(\tilde{y} \mid \mathbf{y})$ via simulation:

This method highlights the power of simulation-based inference.

1. Generate (say) 10,000 simulated parameter values from the posterior distribution. Call these $\theta^{(1)}, \dots, \theta^{(10000)}$.
2. For each $j = 1, \dots, 10000$ separately, generate a value of $\tilde{y}^{(j)}$ from the probability model with parameter value $\theta^{(j)}$.
3. Summarize features of the 10,000 values $\tilde{y}^{(1)}, \dots, \tilde{y}^{(10000)}$ for predictions.

Example with 30-day mortality:

The posterior distribution is $\theta \sim \text{Be}(2, 5)$. Also, $\Pr(\tilde{y} = 1 \mid \theta) = \theta$.

```
> theta = rbeta(10000, 2, 5) # 10,000 values from Be(2,5)
> head(theta)
[1] 0.1979491 0.4944512 0.3124196 0.4441554 0.3895090 0.2404773
> y = rbinom(10000, 1, theta) # generate 10,000 binary values
>                                     # with different probabilities
> head(y)
[1] 0 0 0 1 0 1
> table(y)/10000
y
  0    1
0.7114 0.2886
```

Deeper discussion on prior distributions:

A main difference operationally between Frequentist methods and Bayesian methods is the incorporation of a prior distribution.

What are the guidelines for choosing a prior distribution?

Two types of prior distributions: Informative versus non-informative

Informative prior distributions: The statistician uses his/her knowledge about the substantive problem, along with elicited expert opinion if possible, to construct a prior distribution that properly reflects prior beliefs about the unknown parameters.

Non-informative prior distributions: The statistician chooses a distribution in an attempt to be objective, acting as though no prior knowledge about the parameters exists before observing the data.

When one can construct a defensible informative prior distribution, this is the best and most scientific approach.

Criticisms of assuming an informative prior distribution:

1. Two different Bayesians could end up using two different informative prior distributions, and would therefore obtain two different posterior distributions.
2. Some argue informative prior distributions are not “objective” or “scientific.”

Noninformative prior distributions:

Also termed “vague,” “diffuse,” and “objective.”

It is often desirable to have prior distributions that

- formally express ignorance, and
- can be viewed as default choices when no prior knowledge is available.

Main desired property of choosing an objective prior distribution:

- Assigns “equal probability” (or at least approximately equal) to all values of the parameters.

That is, the prior density should be fairly flat, and that the likelihood should overwhelm the prior density with even small amounts of data.

Problems with objective prior distributions:

1. For most problems, there is no unique non-informative prior distribution.
2. Any method for constructing a non-informative prior distribution should be invariant to the scale of the parameter. This is a difficult property to satisfy in practice.
3. Common methods for constructing non-informative prior distributions result in “improper” probability distributions.

So being Bayesian is not necessarily the perfect solution.

Back to Monte Carlo simulation:

- When posterior densities become more complex, they become very difficult to summarize.

- If the posterior density is non-standard, then straight Monte Carlo simulation can be difficult as well.

In the late 1980s, a few statisticians (Alan Gelfand, Adrian Smith, and others) stumbled on an idea that revolutionized Bayesian statistics.

Main innovation:

Rather than directly simulate from $p(\theta|y)$, simulate a random walk in the space of θ which converges to $p(\theta|y)$.

Markov chain Monte Carlo (MCMC) simulation:

- The key idea is to create a Markov chain whose stationary distribution is $p(\theta|y)$. Values are then computer-simulated from the Markov chain.
- Once the Markov chain is run long enough, simulated values can be treated as coming from the posterior distribution.

There are actually many ways to set up such a Markov chain. The most straightforward way is to implement a **Gibbs sampler**.

Implementing an MCMC sampler to obtain simulated parameter values:

1. Run several parallel Gibbs samplers with different starting values (preferably widely-dispersed)
2. Simulate values from the Markov chains for a “burn-in” period (before the Markov chains have converged to the stationary distribution), and discard the burn-in simulations
3. Save simulated values after burn-in period. These will be the simulated values on which to perform inferential summaries.

If I were you, I would be left with a big headache because a lot of work still seems necessary to implement a Gibbs sampler.

- Still need to determine analytically the probability distributions for each step in a Gibbs sampler
- Still need to simulate from the probability distributions
- Still need to write code that performs the Markov chain simulation, checks for convergence, etc.

Fortunately, many packages now exist that do the hard work for you.

The two most well-known packages for implementing Bayesian analyses using MCMC are

- WinBUGS and its variants (BUGS = **B**ayesian inference **U**sing **G**ibbs **S**ampling)
- Stan, which is a lot like WinBUGS, but uses (by default) a different type of Markov chain sampler

These programs actually figure out the conditional posterior distributions and automatically sample from them!

That's right – they automatically determine and implement the Markov chain!

All you need to do typically is

- specify the probability model for data and the prior distribution
- specify logistical issues concerning MCMC simulation (starting values, burn-in period, how many iterations to run after burn-in)

Aside from the generic programs like WinBUGS and Stan, plenty of programs exists to fit specific common models.

Example: The R library `MCMCpack` implements a variety of common models that are fit via MCMC simulation from the posterior distribution.

Application: Hierarchical linear models

A class of models that is particularly well-served by the Bayesian framework is hierarchical models.

Typical setup: Observe response data and predictor variables as in the usual regression setting.

Suppose now that the samples are clustered by a grouping variable.

Examples:

- Obtain a sample of patients seeing their doctor for chest pain. Want to measure the number of days until pain goes away as a function of patient characteristics and treatment given.
The data may be clustered by clinic, and we might expect that the effects of the patient characteristics differ by clinic.
- Obtain a sample of children taking a standardized achievement test. Want to measure the relationship between test score and background variables about each child.
The data may be clustered by school, and we might expect that the relationship between child characteristics and test score might depend on the school.

Example to implement: Reaction times and sleep deprivation

A study published in a 2003 issue of the *Journal of Sleep Research* measured the reaction time per day for 18 subjects in a sleep deprivation study.


On day 0, each subject had a normal amount of sleep. Starting that night, each subject was restricted to 3 hours of sleep.

The response represents the average reaction time on a series of tests given to each subject across 10 days.


Consider the following two extremes:

- For each subject, fit a separate least-squares regression.
 - Good: Acknowledges that the relationship between x and y may differ within each group.
 - Bad: Very little data on which to estimate effects.
- Pool all the data together and fit a single least-squares regression.
 - Good: Makes full use of the entire data.
 - Bad: Does not recognize differences in effects within groups.

Sleep deprivation – reaction times with least-squares fits



Sleep deprivation – pooling data across all subjects



Hierarchical modeling: A compromise between separate regressions and one overall regression

Suppose for group $g = 1, \dots, G$,

$$y_{ig} = \mathbf{x}'_{ig}\boldsymbol{\beta}_g + \varepsilon_{ig} = \beta_{0g} + \beta_{1g}x_{i1g} + \dots + \beta_{Jg}x_{iJg} + \varepsilon_{ig}$$

with $\varepsilon_{ig} \sim N(0, \sigma^2)$.

Notice that the above model assumes a different $\boldsymbol{\beta}_g$ for each group g .

Now introduce the model component

$$\boldsymbol{\beta}_g \sim N(\boldsymbol{\mu}_\beta, \boldsymbol{\Sigma}_\beta).$$



We will also assume a prior distribution on $\boldsymbol{\mu}_\beta$ and $\boldsymbol{\Sigma}_\beta$.

This normal “random effects” distribution $N(\boldsymbol{\mu}_\beta, \boldsymbol{\Sigma}_\beta)$ has an important role:

- The $\boldsymbol{\beta}_g$ are assumed distinct, but come from the same “family” (i.e., $N(\boldsymbol{\mu}_\beta, \boldsymbol{\Sigma}_\beta)$).
- The data across all groups inform the values of $\boldsymbol{\mu}_\beta$ and $\boldsymbol{\Sigma}_\beta$, which in turn means that the $\boldsymbol{\beta}_g$ are informed from other groups besides the data in group g . Sometimes called “partial pooling” of data across groups.
- Can think of the normal distribution acting like a rubber band around the $\boldsymbol{\beta}_g$. They can vary, but the random effects distribution keeps them from being too far apart.
- The random effects distribution “shrinks” the $\boldsymbol{\beta}_g$ to a common population mean.
- Shrinkage tends to be particularly noticeable when some groups have small numbers of observations relative to others.

Can use the `MCMChregress` function within the `MCMCpack` R library.

```
sleep.mcmc = MCMChregress(fixed = Reaction ~ Days,
  random = ~ Days,
  group = "Subject",
  data=sleepstudy,
  burnin = 5000, mcmc = 10000, # burn-in and total iterations
  thin = 1, verbose = 1,
  beta.start = NA, sigma2.start = NA, Vb.start = NA,
  mubeta = c(250,10), # prior mean of beta_0 and beta_1
  Vbeta = 10000,
  r = 3, R = diag(c(1, 0.1)), nu = 0.001, delta = 0.001)
```



Final thoughts about Bayesian statistics:

- Precise modeling leads to statistical inferences through a leak-proof route
- Given a probability model, Bayesian analysis makes full use of all the data
- Statistical inferences that are unacceptable must come from inappropriate modeling assumptions, not a problem in the underlying inferential mechanism
- Less attention need to be given to mathematical convenience of models, and more attention on scientific merit
- Awkward problems that Frequentists face (e.g., choice of estimators, adjustments for certain types of data) do not arise for Bayesians
- Modern computational methods (MCMC sampling from the posterior distribution) enables fitting even complex data models.